

Task-Driven Detection of Distribution Shifts with Statistical Guarantees for Robot Learning


Authors: A. Farid, S. Veer, D. Pachisia, A. Majumdar

Charis Stamouli

April 19th, 2022



Motivation



Objectives:

- ▶ OOD detection with guaranteed confidence bounds

Objectives:

- ▶ OOD detection with guaranteed confidence bounds
- ▶ Task-driven approach (sensitive only to changes that affect the robot's performance)

Objectives:

- ▶ OOD detection with guaranteed confidence bounds
- ▶ Task-driven approach (sensitive only to changes that affect the robot's performance)

Tools:

- ▶ Probably Approximately Correct (PAC)-Bayes theory

Objectives:

- ▶ OOD detection with guaranteed confidence bounds
- ▶ Task-driven approach (sensitive only to changes that affect the robot's performance)

Tools:

- ▶ Probably Approximately Correct (PAC)-Bayes theory
- ▶ Statistical techniques based on p-values and concentration inequalities

- ▶ Train control policies with a guaranteed performance bound on the training distribution

Contributions

- ▶ Train control policies with a guaranteed performance bound on the training distribution
- ▶ Develop two different task-driven OOD detectors

- ▶ Train control policies with a guaranteed performance bound on the training distribution
- ▶ Develop two different task-driven OOD detectors
- ▶ Perform OOD and WD detection with guaranteed confidence bounds

- ▶ Train control policies with a guaranteed performance bound on the training distribution
- ▶ Develop two different task-driven OOD detectors
- ▶ Perform OOD and WD detection with guaranteed confidence bounds
 - ⇒ Statistical guarantees on both the false positive and false negative rates of the detectors

- ▶ Train control policies with a guaranteed performance bound on the training distribution
- ▶ Develop two different task-driven OOD detectors
- ▶ Perform OOD and WD detection with guaranteed confidence bounds
 - ⇒ Statistical guarantees on both the false positive and false negative rates of the detectors
- ▶ Demonstrate the benefits of the approach in simulation and hardware

Problem formulation

- ▶ Robot dynamics:

$$s_{t+1} = f_E(s_t, a_t)$$

Problem formulation

- ▶ Robot dynamics:

$$s_{t+1} = f_E(s_t, a_t)$$

- ▶ Cost under policy π :

$$C_E(\pi)$$

Problem formulation

- ▶ Robot dynamics:

$$s_{t+1} = f_E(s_t, a_t)$$

- ▶ Cost under policy π :

$$C_E(\pi)$$

- ▶ Training set:

$$S = \{E_1, \dots, E_m\} \sim \mathcal{D}^m$$

Problem formulation

- ▶ Robot dynamics:

$$s_{t+1} = f_E(s_t, a_t)$$

- ▶ Cost under policy π :

$$C_E(\pi)$$

- ▶ Training set:

$$S = \{E_1, \dots, E_m\} \sim \mathcal{D}^m$$

- ▶ Test set:

$$S' = \{E'_1, \dots, E'_n\} \sim \mathcal{D}'^n$$

Problem formulation

- ▶ Robot dynamics:

$$s_{t+1} = f_E(s_t, a_t)$$

- ▶ Cost under policy π :

$$C_E(\pi)$$

- ▶ Training set:

$$S = \{E_1, \dots, E_m\} \sim \mathcal{D}^m$$

- ▶ Test set:

$$S' = \{E'_1, \dots, E'_n\} \sim \mathcal{D}'^n$$

**After deploying the robot in a few environments in S' ,
can we detect if \mathcal{D}' is different from \mathcal{D} ?**

- ▶ OOD detection if:

$$C_{\mathcal{D}'}(\pi) := \mathbb{E}_{E' \sim \mathcal{D}'} C_{E'}(\pi) > C_{\mathcal{D}}(\pi) := \mathbb{E}_{E \sim \mathcal{D}} C_E(\pi)$$

- ▶ OOD detection if:

$$C_{\mathcal{D}'}(\pi) := \mathbb{E}_{E' \sim \mathcal{D}'} C_{E'}(\pi) > C_{\mathcal{D}}(\pi) := \mathbb{E}_{E \sim \mathcal{D}} C_E(\pi)$$

- ▶ WD detection if:

$$C_{\mathcal{D}'}(\pi) \leq C_{\mathcal{D}}(\pi)$$

- ▶ OOD detection if:

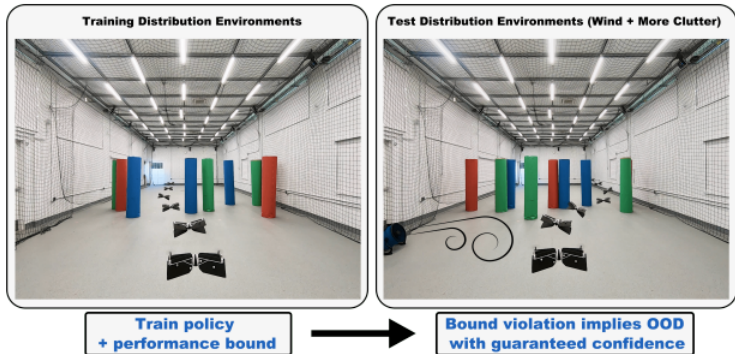
$$C_{\mathcal{D}'}(\pi) := \mathbb{E}_{E' \sim \mathcal{D}'} C_{E'}(\pi) > C_{\mathcal{D}}(\pi) := \mathbb{E}_{E \sim \mathcal{D}} C_E(\pi)$$

- ▶ WD detection if:

$$C_{\mathcal{D}'}(\pi) \leq C_{\mathcal{D}}(\pi)$$

- ▶ Detection: insensitive to changes in the environment distribution that do not adversely impact the robot's performance

Approach



Policy training via PAC-Bayes bounds

- ▶ Goal: Learn a policy π with a guaranteed bound on the expected cost $C_{\mathcal{D}}(\pi)$

Policy training via PAC-Bayes bounds

- ▶ Goal: Learn a policy π with a guaranteed bound on the expected cost $C_{\mathcal{D}}(\pi)$
- ▶ PAC-Bayes:
 - ▶ Before observing any data choose a **prior distribution** P_0 over the policy space Π

Policy training via PAC-Bayes bounds

- ▶ Goal: Learn a policy π with a guaranteed bound on the expected cost $C_{\mathcal{D}}(\pi)$
- ▶ PAC-Bayes:
 - ▶ Before observing any data choose a **prior distribution** P_0 over the policy space Π
 - ▶ Obtain a training dataset S and choose a **posterior distribution** P over the policy space Π

Policy training via PAC-Bayes bounds

- ▶ Goal: Learn a policy π with a guaranteed bound on the expected cost $C_{\mathcal{D}}(\pi)$
- ▶ PAC-Bayes:
 - ▶ Before observing any data choose a **prior distribution** P_0 over the policy space Π
 - ▶ Obtain a training dataset S and choose a **posterior distribution** P over the policy space Π
- ▶ Training cost:

$$C_S(\pi) = \frac{1}{m} \sum_{E \in S} C_E(\pi)$$

Policy training via PAC-Bayes bounds

Theorem 1. *For any distribution \mathcal{D} , prior distribution P_0 , $\delta \in (0, 1)$, cost bounded in $[0, 1]$, $m \geq 8$, and deterministic algorithm A which outputs the posterior distribution P , we have the following:*

$$\mathbb{P}_{(S, \pi) \sim (\mathcal{D}^m \times P)} [C_{\mathcal{D}}(\pi) \leq \bar{C}_{\delta}(\pi, S)] \geq 1 - \delta,$$

where $\bar{C}_{\delta}(\pi, S) := C_S(\pi) + \sqrt{R}$, $R := \left(D_2(P||P_0) + \ln \left(\frac{2\sqrt{m}}{(\delta/2)^3} \right) \right) / (2m)$,

and $D_2(P||P_0) := \ln \left(\mathbb{E}_{\pi \sim P_0} \left[\left(\frac{P(\pi)}{P_0(\pi)} \right)^2 \right] \right)$ is the Rényi divergence.

Policy training via PAC-Bayes bounds

Theorem 1. *For any distribution \mathcal{D} , prior distribution P_0 , $\delta \in (0, 1)$, cost bounded in $[0, 1]$, $m \geq 8$, and deterministic algorithm A which outputs the posterior distribution P , we have the following:*

$$\mathbb{P}_{(S, \pi) \sim (\mathcal{D}^m \times P)} [C_{\mathcal{D}}(\pi) \leq \bar{C}_{\delta}(\pi, S)] \geq 1 - \delta,$$

where $\bar{C}_{\delta}(\pi, S) := C_S(\pi) + \sqrt{R}$, $R := \left(D_2(P||P_0) + \ln \left(\frac{2\sqrt{m}}{(\delta/2)^3} \right) \right) / (2m)$,

and $D_2(P||P_0) := \ln \left(\mathbb{E}_{\pi \sim P_0} \left[\left(\frac{P(\pi)}{P_0(\pi)} \right)^2 \right] \right)$ is the Rényi divergence.

Corollary 1. *Let the assumptions of Theorem 1 hold. Then:*

$$\mathbb{P}_{(S, \pi) \sim (\mathcal{D}^m \times P)} [C_{\mathcal{D}}(\pi) \geq \underline{C}_{\delta}(\pi, S)] \geq 1 - \delta,$$

where $\underline{C}_{\delta}(\pi, S) := C_S(\pi) - \sqrt{R}$.

Policy training via PAC-Bayes bounds

- ▶ PAC-Bayes bounds: Strong bounds for policy learning

Policy training via PAC-Bayes bounds

- ▶ PAC-Bayes bounds: Strong bounds for policy learning
- ▶ Policy training: Search for a posterior P to minimize the upper bound $\bar{C}_\delta(\pi, S)$

Policy training via PAC-Bayes bounds

- ▶ PAC-Bayes bounds: Strong bounds for policy learning
- ▶ Policy training: Search for a posterior P to minimize the upper bound $\bar{C}_\delta(\pi, S)$
- ▶ Training methods:
 - ▶ Backpropagation
 - ▶ Blackbox optimization (evolutionary strategies)

- ▶ Key idea:

Violation of the upper bound $\bar{C}_\delta(\pi, S)$



Test environment is OOD

Task-driven OOD detection with statistical guarantees

- ▶ Key idea:

Violation of the upper bound $\bar{C}_\delta(\pi, S)$



Test environment is OOD

- ▶ Methods:

1. Hypothesis testing via p-value

Task-driven OOD detection with statistical guarantees

► Key idea:

Violation of the upper bound $\bar{C}_\delta(\pi, S)$



Test environment is OOD

► Methods:

1. Hypothesis testing via p-value
2. Confidence interval on the difference in expected train and test costs

Method 1: Hypothesis testing

- ▶ Goal: Detect if the test dataset is **OOD**, **WD** or **Unknown**

Method 1: Hypothesis testing

- ▶ Goal: Detect if the test dataset is **OOD**, **WD** or **Unknown**
- ▶ Define p-values for:
 - ▶ OOD detection:
$$p_O(S') := \mathbb{P}_{\hat{S}' \sim \mathcal{D}'^n} [C_{\hat{S}'}(\pi) \geq C_{S'}(\pi) | C_{\mathcal{D}'}(\pi) \leq C_{\mathcal{D}}(\pi)]$$

Method 1: Hypothesis testing

- ▶ Goal: Detect if the test dataset is **OOD**, **WD** or **Unknown**
- ▶ Define p-values for:
 - ▶ OOD detection:
$$p_O(S') := \mathbb{P}_{\hat{S}' \sim \mathcal{D}'^n} [C_{\hat{S}'}(\pi) \geq C_{S'}(\pi) | C_{\mathcal{D}'}(\pi) \leq C_{\mathcal{D}}(\pi)]$$
 - ▶ WD detection:
$$p_W(S') := \mathbb{P}_{\hat{S}' \sim \mathcal{D}'^n} [C_{\hat{S}'}(\pi) \leq C_{S'}(\pi) | C_{\mathcal{D}'}(\pi) > C_{\mathcal{D}}(\pi)]$$

Method 1: Hypothesis testing

- ▶ Goal: Detect if the test dataset is **OOD**, **WD** or **Unknown**
- ▶ Define p-values for:
 - ▶ OOD detection:
$$p_O(S') := \mathbb{P}_{\hat{S}' \sim \mathcal{D}'^n} [C_{\hat{S}'}(\pi) \geq C_{S'}(\pi) | C_{\mathcal{D}'}(\pi) \leq C_{\mathcal{D}}(\pi)]$$
 - ▶ WD detection:
$$p_W(S') := \mathbb{P}_{\hat{S}' \sim \mathcal{D}'^n} [C_{\hat{S}'}(\pi) \leq C_{S'}(\pi) | C_{\mathcal{D}'}(\pi) > C_{\mathcal{D}}(\pi)]$$
- ▶ Perform two hypothesis tests:
 - ▶ H_0 : WD and H_1 : OOD

Method 1: Hypothesis testing

- ▶ Goal: Detect if the test dataset is **OOD**, **WD** or **Unknown**
- ▶ Define p-values for:
 - ▶ OOD detection:
$$p_O(S') := \mathbb{P}_{\hat{S}' \sim \mathcal{D}'^n} [C_{\hat{S}'}(\pi) \geq C_{S'}(\pi) | C_{\mathcal{D}'}(\pi) \leq C_{\mathcal{D}}(\pi)]$$
 - ▶ WD detection:
$$p_W(S') := \mathbb{P}_{\hat{S}' \sim \mathcal{D}'^n} [C_{\hat{S}'}(\pi) \leq C_{S'}(\pi) | C_{\mathcal{D}'}(\pi) > C_{\mathcal{D}}(\pi)]$$
- ▶ Perform two hypothesis tests:
 - ▶ H_0 : WD and H_1 : OOD
$$p_O(S') < \alpha_O \implies \mathcal{D}': \text{OOD}$$

Method 1: Hypothesis testing

- ▶ Goal: Detect if the test dataset is **OOD**, **WD** or **Unknown**
- ▶ Define p-values for:
 - ▶ OOD detection:
$$p_O(S') := \mathbb{P}_{\hat{S}' \sim \mathcal{D}'^n} [C_{\hat{S}'}(\pi) \geq C_{S'}(\pi) | C_{\mathcal{D}'}(\pi) \leq C_{\mathcal{D}}(\pi)]$$
 - ▶ WD detection:
$$p_W(S') := \mathbb{P}_{\hat{S}' \sim \mathcal{D}'^n} [C_{\hat{S}'}(\pi) \leq C_{S'}(\pi) | C_{\mathcal{D}'}(\pi) > C_{\mathcal{D}}(\pi)]$$
- ▶ Perform two hypothesis tests:
 - ▶ H_0 : WD and H_1 : OOD
$$p_O(S') < \alpha_O \implies \mathcal{D}': \text{OOD}$$
 - ▶ H_0 : OOD and H_1 : WD
$$p_W(S') < \alpha_W \implies \mathcal{D}': \text{WD}$$

Method 1: Hypothesis testing

► Challenge:

$\mathcal{D}, \mathcal{D}'$ are unknown



Cannot compute the p-values

Method 1: Hypothesis testing

- Challenge:

$\mathcal{D}, \mathcal{D}'$ are unknown



Cannot compute the p-values

- Solution: Compute upper bounds for the p-values:

Method 1: Hypothesis testing

- Challenge:

$\mathcal{D}, \mathcal{D}'$ are unknown



Cannot compute the p-values

- Solution: Compute upper bounds for the p-values:

$$\mathbb{P}_{(S,\pi) \sim (\mathcal{D}^m \times P)} [p_O(S') \leq \exp(-2n \max\{C_{S'}(\pi) - \bar{C}_{\delta_O}(\pi, S), 0\}^2) \geq 1 - \delta_O$$

$$\mathbb{P}_{(S,\pi) \sim (\mathcal{D}^m \times P)} [p_W(S') \leq \exp(-2n \max\{\underline{C}_{\delta_W}(\pi, S) - C_{S'}(\pi), 0\}^2) \geq 1 - \delta_W$$

Method 1: Hypothesis testing

- Challenge:

$\mathcal{D}, \mathcal{D}'$ are unknown



Cannot compute the p-values

- Solution: Compute upper bounds for the p-values:

$$\mathbb{P}_{(S,\pi) \sim (\mathcal{D}^m \times P)} [p_O(S') \leq \exp(-2n \max\{C_{S'}(\pi) - \bar{C}_{\delta_O}(\pi, S), 0\}^2)] \geq 1 - \delta_O$$

$$\mathbb{P}_{(S,\pi) \sim (\mathcal{D}^m \times P)} [p_W(S') \leq \exp(-2n \max\{\underline{C}_{\delta_W}(\pi, S) - C_{S'}(\pi), 0\}^2)] \geq 1 - \delta_W$$

- Compare the upper bounds with the significance levels

Method 1: Hypothesis testing

Algorithm 1 OOD/WD Detection using Hypothesis Testing

Input: $\delta_O, \delta_W, \alpha_O, \alpha_W \in (0, 1)$.

Input: PAC-Bayes Bounds: $\overline{C}_{\delta_O}(\pi, S), \underline{C}_{\delta_W}(\pi, S)$.

Input: Test dataset $S' \sim \mathcal{D}'^n$ and policy $\pi \sim P$.

Output: OOD, WD, and Unknown

$C_{S'}(\pi) \leftarrow \frac{1}{n} \sum_{E \in S'} C_E(\pi)$

$\bar{\tau} \leftarrow \max\{C_{S'}(\pi) - \overline{C}_{\delta_O}(\pi, S), 0\}$

$\underline{\tau} \leftarrow \max\{\underline{C}_{\delta_W}(\pi, S) - C_{S'}(\pi), 0\}$

if $\exp(-2n\bar{\tau}^2) \leq \alpha_O$ **then**

 OOD \leftarrow True

end if

if $\exp(-2n\underline{\tau}^2) \leq \alpha_W$ **then**

 WD \leftarrow True

end if

if $\exp(-2n\bar{\tau}^2) > \alpha_W$ and $\exp(-2n\underline{\tau}^2) > \alpha_O$ **then**

 Unknown \leftarrow True

end if

Method 1: Hypothesis testing

Algorithm 1 OOD/WD Detection using Hypothesis Testing

Input: $\delta_O, \delta_W, \alpha_O, \alpha_W \in (0, 1)$.
Input: PAC-Bayes Bounds: $\overline{C}_{\delta_O}(\pi, S), \underline{C}_{\delta_W}(\pi, S)$.
Input: Test dataset $S' \sim \mathcal{D}'^n$ and policy $\pi \sim P$.
Output: OOD, WD, and Unknown
 $C_{S'}(\pi) \leftarrow \frac{1}{n} \sum_{E \in S'} C_E(\pi)$
 $\bar{\tau} \leftarrow \max\{C_{S'}(\pi) - \overline{C}_{\delta_O}(\pi, S), 0\}$
 $\underline{\tau} \leftarrow \max\{\underline{C}_{\delta_W}(\pi, S) - C_{S'}(\pi), 0\}$
if $\exp(-2n\bar{\tau}^2) \leq \alpha_O$ **then**
 OOD \leftarrow True
end if
if $\exp(-2n\underline{\tau}^2) \leq \alpha_W$ **then**
 WD \leftarrow True
end if
if $\exp(-2n\bar{\tau}^2) > \alpha_W$ and $\exp(-2n\underline{\tau}^2) > \alpha_O$ **then**
 Unknown \leftarrow True
end if

- Consistency: Algorithm 1 returns **only one of the three possible outcomes** (OOD, WD or Unknown)

Method 2: Confidence interval overlap

- ▶ Goal: Detect if the test dataset is **OOD**, **WD** or **Unknown**

Method 2: Confidence interval overlap

- Goal: Detect if the test dataset is **OOD**, **WD** or **Unknown**

Theorem 3. Let \mathcal{D} be the training distribution, \mathcal{D}' be the test distribution, and P be the posterior distribution on the policy space obtained through training. Let $\delta_O, \delta'_O \in (0, 1)$ such that $\delta_O + \delta'_O < 1$, $\gamma_O := \sqrt{\frac{\ln(1/\delta'_O)}{2n}}$, and $\Delta C_O := C_{S'}(\pi) - \gamma_O - \bar{C}_{\delta_O}(\pi, S)$. Similarly, let $\delta_W, \delta'_W \in (0, 1)$ such that $\delta_W + \delta'_W < 1$, $\gamma_W := \sqrt{\frac{\ln(1/\delta'_W)}{2n}}$, and $\Delta C_W := \underline{C}_{\delta_W}(\pi, S) - C_{S'}(\pi) - \gamma_W$. Then,

$$\mathbb{P}_{(S, \pi, S') \sim (\mathcal{D}^m \times P \times \mathcal{D}'^n)} [C_{\mathcal{D}'}(\pi) - C_{\mathcal{D}}(\pi) \geq \Delta C_O] \geq 1 - \delta_O - \delta'_O,$$

$$\mathbb{P}_{(S, \pi, S') \sim (\mathcal{D}^m \times P \times \mathcal{D}'^n)} [C_{\mathcal{D}}(\pi) - C_{\mathcal{D}'}(\pi) \geq \Delta C_W] \geq 1 - \delta_W - \delta'_W.$$

Method 2: Confidence interval overlap

- ▶ ΔC_O : high-confidence lower bound for $C_{\mathcal{D}'}(\pi) - C_{\mathcal{D}}(\pi)$

Method 2: Confidence interval overlap

- ▶ ΔC_O : high-confidence lower bound for $C_{\mathcal{D}'}(\pi) - C_{\mathcal{D}}(\pi)$

$$\Delta C_O > 0 \implies \mathcal{D}' : \text{OOD}$$

Method 2: Confidence interval overlap

- ▶ ΔC_O : high-confidence lower bound for $C_{\mathcal{D}'}(\pi) - C_{\mathcal{D}}(\pi)$

$$\Delta C_O > 0 \implies \mathcal{D}' : \text{OOD}$$

- ▶ ΔC_W : high-confidence lower bound for $C_{\mathcal{D}}(\pi) - C_{\mathcal{D}'}(\pi)$

Method 2: Confidence interval overlap

- ▶ ΔC_O : high-confidence lower bound for $C_{\mathcal{D}'}(\pi) - C_{\mathcal{D}}(\pi)$

$$\Delta C_O > 0 \implies \mathcal{D}' : \text{OOD}$$

- ▶ ΔC_W : high-confidence lower bound for $C_{\mathcal{D}}(\pi) - C_{\mathcal{D}'}(\pi)$

$$\Delta C_W \geq 0 \implies \mathcal{D}' : \text{WD}$$

Method 2: Confidence interval overlap

- ▶ ΔC_O : high-confidence lower bound for $C_{\mathcal{D}'}(\pi) - C_{\mathcal{D}}(\pi)$

$$\Delta C_O > 0 \implies \mathcal{D}' : \text{OOD}$$

- ▶ ΔC_W : high-confidence lower bound for $C_{\mathcal{D}}(\pi) - C_{\mathcal{D}'}(\pi)$

$$\Delta C_W \geq 0 \implies \mathcal{D}' : \text{WD}$$

- ▶ $\Delta C_O \leq 0$ and $\Delta C_W < 0 \implies \mathcal{D}' : \text{Unknown}$

Method 2: Confidence interval overlap

Algorithm 2 OOD/WD Detection using Confidence Intervals

Input: $\delta_O, \delta'_O \in (0, 1)$ with desired **maximum** false positive rate $\delta_O + \delta'_O < 1$.
Input: $\delta_W, \delta'_W \in (0, 1)$ with desired **maximum** false negative rate, $\delta_W + \delta'_W < 1$.
Input: PAC-Bayes Bounds: $\overline{C}_{\delta_O}(\pi, S)$, $\underline{C}_{\delta_W}(\pi, S)$.
Input: Test dataset $S' \sim \mathcal{D}^n$ and policy $\pi \sim P$.
Output: OOD, WD and Unknown.

$$C_{S'}(\pi) \leftarrow \frac{1}{n} \sum_{E' \in S'} C_{E'}(\pi)$$
$$\gamma_O \leftarrow \sqrt{\frac{\ln(1/\delta'_O)}{2n}}$$
$$\gamma_W \leftarrow \sqrt{\frac{\ln(1/\delta'_W)}{2n}}$$
$$\Delta C_O \leftarrow C_{S'}(\pi) - \gamma_O - \overline{C}_{\delta_O}(\pi, S)$$
$$\Delta C_W \leftarrow \underline{C}_{\delta_W}(\pi, S) - C_{S'}(\pi) - \gamma_W$$

if $\Delta C_O > 0$ **then**
 OOD \leftarrow True
end if
if $\Delta C_W \geq 0$ **then**
 WD \leftarrow True
end if
if $\Delta C_O \leq 0$ and $\Delta C_W < 0$ **then**
 Unknown \leftarrow True
end if

Method 2: Confidence interval overlap

Algorithm 2 OOD/WD Detection using Confidence Intervals

Input: $\delta_O, \delta'_O \in (0, 1)$ with desired **maximum** false positive rate $\delta_O + \delta'_O < 1$.
Input: $\delta_W, \delta'_W \in (0, 1)$ with desired **maximum** false negative rate, $\delta_W + \delta'_W < 1$.
Input: PAC-Bayes Bounds: $\overline{C}_{\delta_O}(\pi, S)$, $\underline{C}_{\delta_W}(\pi, S)$.
Input: Test dataset $S' \sim \mathcal{D}^n$ and policy $\pi \sim P$.
Output: OOD, WD and Unknown.
 $C_{S'}(\pi) \leftarrow \frac{1}{n} \sum_{E' \in S'} C_{E'}(\pi)$
 $\gamma_O \leftarrow \sqrt{\frac{\ln(1/\delta'_O)}{2n}}$
 $\gamma_W \leftarrow \sqrt{\frac{\ln(1/\delta'_W)}{2n}}$
 $\Delta C_O \leftarrow C_{S'}(\pi) - \gamma_O - \overline{C}_{\delta_O}(\pi, S)$
 $\Delta C_W \leftarrow \underline{C}_{\delta_W}(\pi, S) - C_{S'}(\pi) - \gamma_W$
if $\Delta C_O > 0$ **then**
 OOD \leftarrow True
end if
if $\Delta C_W \geq 0$ **then**
 WD \leftarrow True
end if
if $\Delta C_O \leq 0$ and $\Delta C_W < 0$ **then**
 Unknown \leftarrow True
end if

- Consistency: Algorithm 2 returns **only one of the three possible outcomes** (OOD, WD or Unknown)

- ▶ Franka Panda arm with an overhead camera

Robotic grasping

- ▶ Franka Panda arm with an overhead camera
- ▶ Train the manipulator to grasp mugs placed in poses drawn from a specific distribution

Robotic grasping

- ▶ Franka Panda arm with an overhead camera
- ▶ Train the manipulator to grasp mugs placed in poses drawn from a specific distribution
- ▶ Control policy: DNN which inputs a depth map of the object and a latent state

Robotic grasping

- ▶ Franka Panda arm with an overhead camera
- ▶ Train the manipulator to grasp mugs placed in poses drawn from a specific distribution
- ▶ Control policy: DNN which inputs a depth map of the object and a latent state
- ▶ Training:
 - ▶ Dataset: Mugs from ShapeNet randomly scaled in all dimensions

Robotic grasping

- ▶ Franka Panda arm with an overhead camera
- ▶ Train the manipulator to grasp mugs placed in poses drawn from a specific distribution
- ▶ Control policy: DNN which inputs a depth map of the object and a latent state
- ▶ Training:
 - ▶ Dataset: Mugs from ShapeNet randomly scaled in all dimensions
 - ▶ Cost: Successful rollout if the robot is able to lift the mug by 10cm and the gripper palm does not contact it

Robotic grasping

- ▶ Franka Panda arm with an overhead camera
- ▶ Train the manipulator to grasp mugs placed in poses drawn from a specific distribution
- ▶ Control policy: DNN which inputs a depth map of the object and a latent state
- ▶ Training:
 - ▶ Dataset: Mugs from ShapeNet randomly scaled in all dimensions
 - ▶ Cost: Successful rollout if the robot is able to lift the mug by 10cm and the gripper palm does not contact it
 - ▶ Process: Sample a policy π from the trained posterior

Robotic grasping: Simulation results


- ▶ Gradually make the distribution on the mug poses more challenging

Robotic grasping: Simulation results

- ▶ Gradually make the distribution on the mug poses more challenging
- ▶ For each distribution:
 - ▶ Sample 20 test datasets
 - ▶ Compute the OOD indicators

Robotic grasping: Simulation results

- ▶ Gradually make the distribution on the mug poses more challenging
- ▶ For each distribution:
 - ▶ Sample 20 test datasets
 - ▶ Compute the OOD indicators




Robotic grasping: Simulation results

- ▶ Change the objects from mugs to bowls



Robotic grasping: Simulation results

- Change the objects from mugs to bowls



Robotic grasping: Simulation results

- Change the objects from mugs to bowls



Robotic grasping: Hardware results

- ▶ Gradually increase the range of location of the mugs

Robotic grasping: Hardware results

- ▶ Gradually increase the range of location of the mugs




- ▶ The proportion of trials that the grasp fails increases

Robotic grasping: Hardware results

- Gradually increase the range of location of the mugs



- The proportion of trials that the grasp fails increases




Robotic grasping: Hardware results

- ▶ Change the objects from mugs to bowls



Robotic grasping: Hardware results

- Change the objects from mugs to bowls



Robotic grasping: Hardware results

- Change the objects from mugs to bowls



- ▶ Parrot Swing drone

Obstacle avoidance with a drone

- ▶ Parrot Swing drone
- ▶ Train the drone to avoid obstacles by the largest possible distance

Obstacle avoidance with a drone

- ▶ Parrot Swing drone
- ▶ Train the drone to avoid obstacles by the largest possible distance
- ▶ Control policy: DNN which inputs a depth image and outputs a set of motion primitives

Obstacle avoidance with a drone

- ▶ Parrot Swing drone
- ▶ Train the drone to avoid obstacles by the largest possible distance
- ▶ Control policy: DNN which inputs a depth image and outputs a set of motion primitives
- ▶ Training:
 - ▶ Dataset: Environments with randomly placed cylindrical obstacles

Obstacle avoidance with a drone

- ▶ Parrot Swing drone
- ▶ Train the drone to avoid obstacles by the largest possible distance
- ▶ Control policy: DNN which inputs a depth image and outputs a set of motion primitives
- ▶ Training:
 - ▶ Dataset: Environments with randomly placed cylindrical obstacles
 - ▶ Cost: $\max \left\{ 0, 1 - \frac{d_{\min}}{d_{thres}} \right\}$

Obstacle avoidance with a drone

- ▶ Parrot Swing drone
- ▶ Train the drone to avoid obstacles by the largest possible distance
- ▶ Control policy: DNN which inputs a depth image and outputs a set of motion primitives
- ▶ Training:
 - ▶ Dataset: Environments with randomly placed cylindrical obstacles
 - ▶ Cost: $\max \left\{ 0, 1 - \frac{d_{\min}}{d_{\text{thres}}} \right\}$
 - ▶ Process: Sample a policy π from the trained posterior

Obstacle avoidance with a drone: Simulation results

- ▶ Generate test datasets of varying difficulty (by changing the number of obstacles and the maximum/minimum gap-size between obstacles)

Obstacle avoidance with a drone: Simulation results



- ▶ Generate test datasets of varying difficulty (by changing the number of obstacles and the maximum/minimum gap-size between obstacles)
- ▶ For each difficulty setting:
 - ▶ Sample 2000 datasets

Obstacle avoidance with a drone: Simulation results

- ▶ Generate test datasets of varying difficulty (by changing the number of obstacles and the maximum/minimum gap-size between obstacles)
- ▶ For each difficulty setting:
 - ▶ Sample 2000 datasets
 - ▶ Evaluate the proportion of datasets that the detectors declare OOD, WD or Unknown

Obstacle avoidance with a drone: Simulation results

- ▶ Generate test datasets of varying difficulty (by changing the number of obstacles and the maximum/minimum gap-size between obstacles)
- ▶ For each difficulty setting:
 - ▶ Sample 2000 datasets
 - ▶ Evaluate the proportion of datasets that the detectors declare OOD, WD or Unknown





Obstacle avoidance with a drone: Simulation results

- ▶ Change the desired maximum false positive and false negative rates

Obstacle avoidance with a drone: Simulation results

- Change the desired maximum false positive and false negative rates




- ▶ Compare the detectors with two baselines:

Obstacle avoidance with a drone: Simulation results

- ▶ Compare the detectors with two baselines:
 - ▶ Maximum softmax probability (MSP)
 - ▶ MaxLogit


Obstacle avoidance with a drone: Simulation results

- ▶ Compare the detectors with two baselines:
 - ▶ Maximum softmax probability (MSP)
 - ▶ MaxLogit



Obstacle avoidance with a drone: Simulation results


- ▶ Compare the detectors with two baselines:
 - ▶ Maximum softmax probability (MSP)
 - ▶ MaxLogit



- ▶ Disadvantages of both baselines:

Obstacle avoidance with a drone: Simulation results


- ▶ Compare the detectors with two baselines:
 - ▶ Maximum softmax probability (MSP)
 - ▶ MaxLogit



- ▶ Disadvantages of both baselines:
 - ▶ Violation of the maximum false positive rate

Obstacle avoidance with a drone: Simulation results

- ▶ Compare the detectors with two baselines:
 - ▶ Maximum softmax probability (MSP)
 - ▶ MaxLogit



- ▶ Disadvantages of both baselines:
 - ▶ Violation of the maximum false positive rate
 - ▶ Triggered by task-irrelevant shifts

Obstacle avoidance with a drone: Hardware results

Deploy the policy trained in simulation on three kinds of environments:

Obstacle avoidance with a drone: Hardware results

Deploy the policy trained in simulation on three kinds of environments:

- ▶ Environments with a smaller number of obstacles (easier environments)

Obstacle avoidance with a drone: Hardware results

Deploy the policy trained in simulation on three kinds of environments:

- ▶ Environments with a smaller number of obstacles (easier environments)
- ▶ Environments with smaller gaps between obstacles (harder environments)

Obstacle avoidance with a drone: Hardware results


Deploy the policy trained in simulation on three kinds of environments:

- ▶ Environments with a smaller number of obstacles (easier environments)
- ▶ Environments with smaller gaps between obstacles (harder environments)
- ▶ Environments with wind generated using a fan

Obstacle avoidance with a drone: Hardware results

Deploy the policy trained in simulation on three kinds of environments:

- ▶ Environments with a smaller number of obstacles (easier environments)
- ▶ Environments with smaller gaps between obstacles (harder environments)
- ▶ Environments with wind generated using a fan



Experiments

- ▶ Videos of the experiments can be found here:
<https://www.youtube.com/watch?v=jKye3A09le0>

Task-driven OOD and WD detection with statistical guarantees:

Task-driven OOD and WD detection with statistical guarantees:

- ▶ Use PAC-Bayes theory to train a policy with a bound on the expected cost on the training distribution

Task-driven OOD and WD detection with statistical guarantees:

- ▶ Use PAC-Bayes theory to train a policy with a bound on the expected cost on the training distribution
- ▶ Perform OOD and WD detection on test environments by checking for violations of the bound using approaches based on both p-values and confidence intervals

Task-driven OOD and WD detection with statistical guarantees:

- ▶ Use PAC-Bayes theory to train a policy with a bound on the expected cost on the training distribution
- ▶ Perform OOD and WD detection on test environments by checking for violations of the bound using approaches based on both p-values and confidence intervals
- ▶ Provide guarantees on the maximum false positive and false negative rate

Task-driven OOD and WD detection with statistical guarantees:

- ▶ Use PAC-Bayes theory to train a policy with a bound on the expected cost on the training distribution
- ▶ Perform OOD and WD detection on test environments by checking for violations of the bound using approaches based on both p-values and confidence intervals
- ▶ Provide guarantees on the maximum false positive and false negative rate
- ▶ Demonstrate the ability of the approaches to perform OOD detection within a handful of trials

Task-driven OOD and WD detection with statistical guarantees:

- ▶ Use PAC-Bayes theory to train a policy with a bound on the expected cost on the training distribution
- ▶ Perform OOD and WD detection on test environments by checking for violations of the bound using approaches based on both p-values and confidence intervals
- ▶ Provide guarantees on the maximum false positive and false negative rate
- ▶ Demonstrate the ability of the approaches to perform OOD detection within a handful of trials
- ▶ Tune the detectors' sensitivity by varying the maximum permissible false positive/negative rate